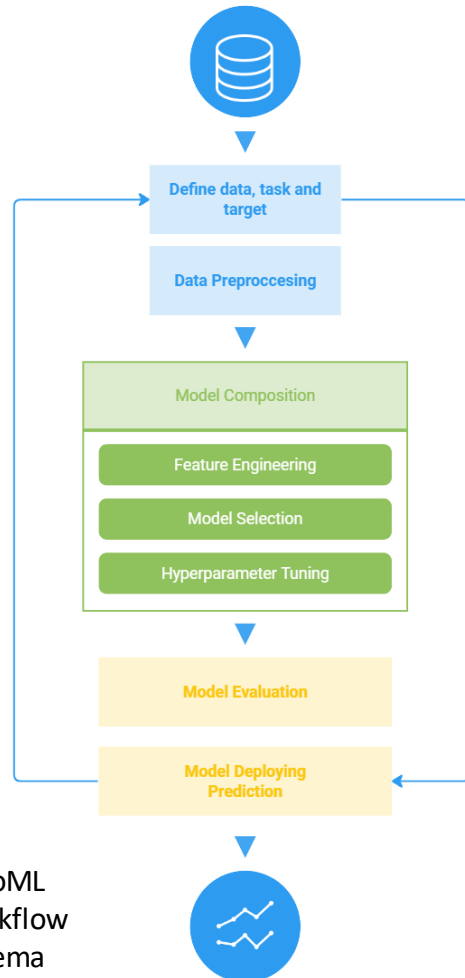


**12th International Young Scientists Conference on Computational Science**



# **Automated Generation of Ensemble Pipelines using Policy-Based Reinforcement Learning method**

**Andrey S. Stebenkov, Nikolay O. Nikitin**  
**ITMO University**



AutoML  
Workflow  
schema

The central concept of **AutoML** is to build an automated workflow that could take raw data as input and produce a prediction automatically.

This automated workflow should automatically do **preprocessing**, **model selection**, **hyperparameter tuning**, and all other stages of the ML process.


The final workflow is called a **Pipeline**.

There are a considerable number of different automated machine learning frameworks. The current widely used approaches are:

- Searching best from predefined models (e.g., AutoSklearn, AutoGluon);
- Constructing "step-by-step" with evolutionary algorithm (e.g., TPOT, FEDOT).

The **paper** introduces an approach to generate ensemble pipelines using **policy-based reinforcement learning**.

# Problem statement



Agent-Environment  
Interaction in AutoML  
schema

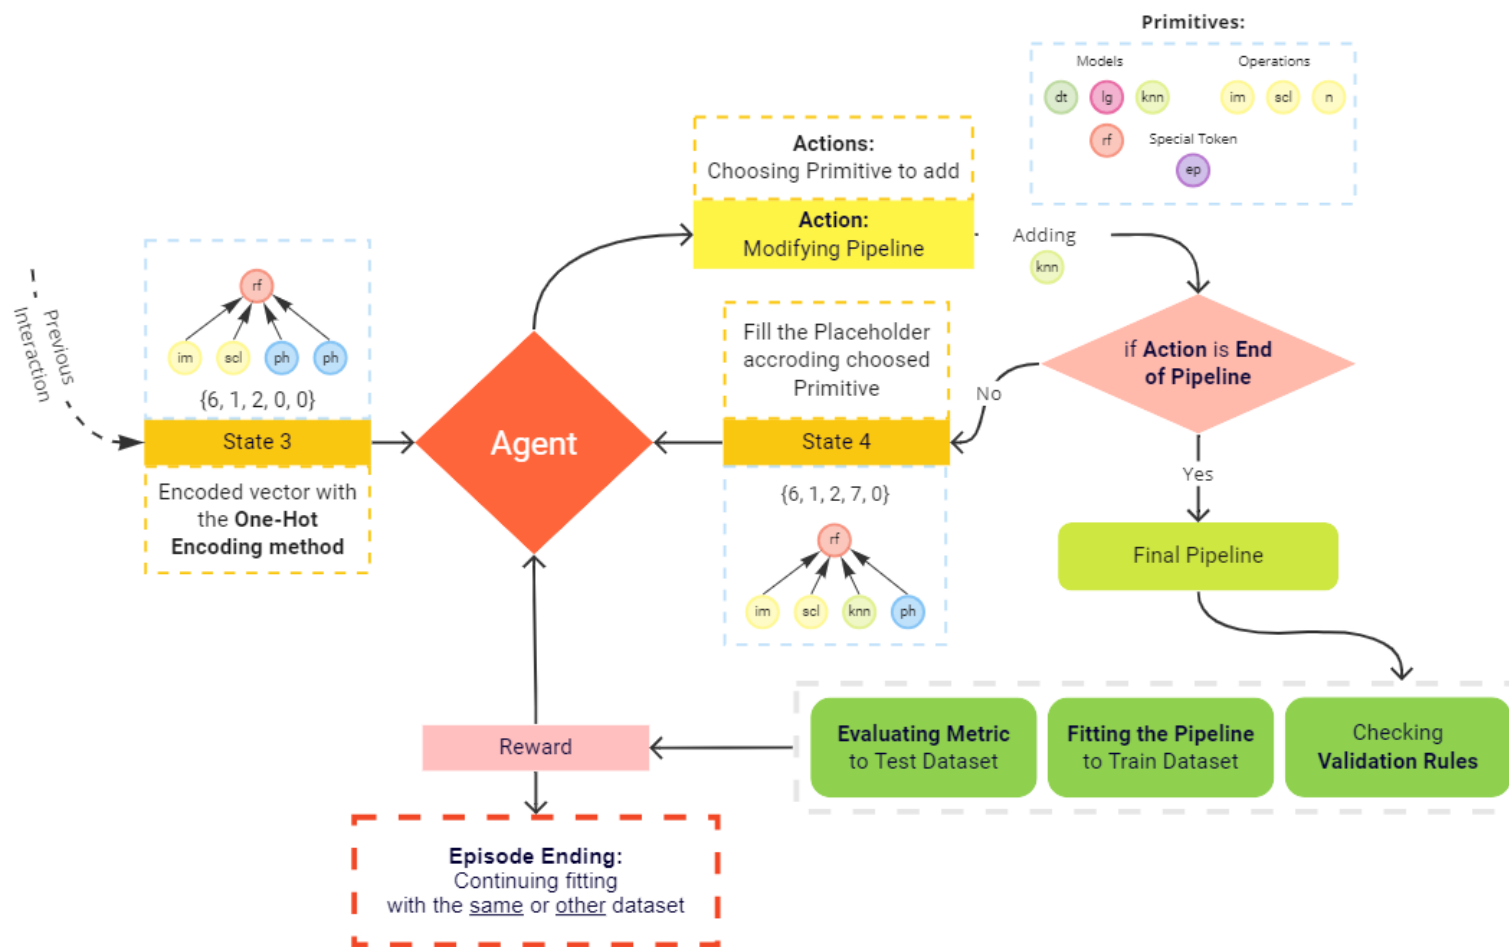
Implementing reinforcement learning to AutoML is not a novel idea. Similar to its predecessors, we suppose that the **pipeline generation problem** can be represented as a **computer game**.

The Agent learns how to **construct step-by-step** the optimal Pipeline by achieving maximum reward using **provided primitives** (models or operations).

**State** — Initial, Intermediate or Final pipeline structure;  
**Action** — Modifying the pipeline (e.g., adding a new model);  
**Reward** — Metric value of the final pipeline;  
**Episode** — Generation one pipeline for the ML task;

# Proposed approach

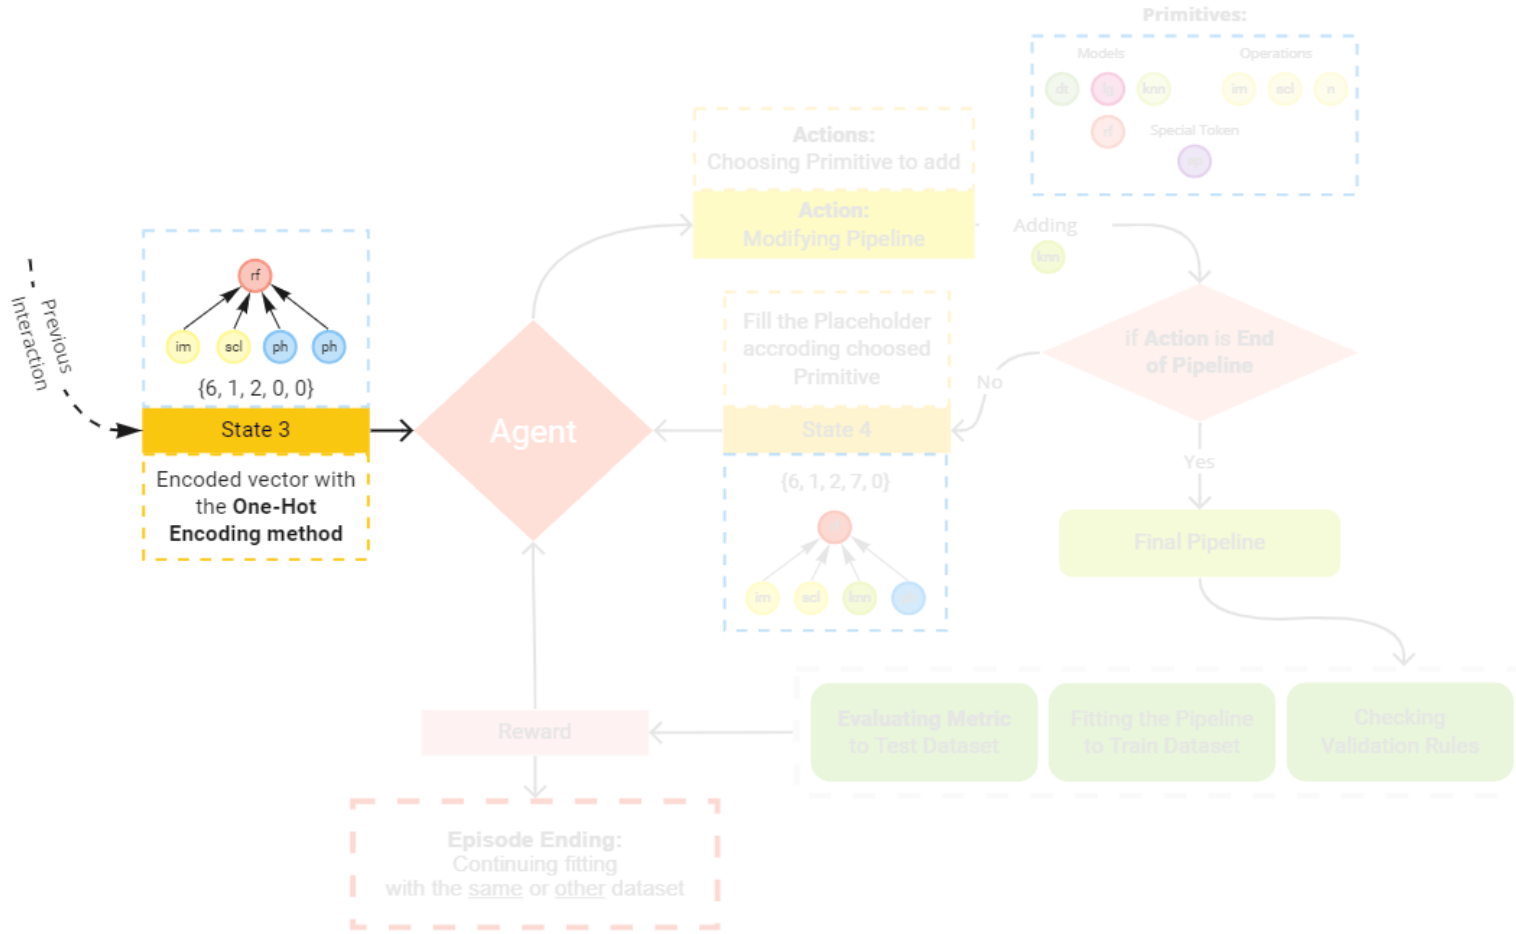
## Full schema of fitting agent




This schema shows the one stage from the entire cycle of Agent's fitting using a proposed approach.

# Proposed approach

## Pipeline Encoding




The stage will start with a pipeline graph into a one-hot encoding vector.



The initial structure of the Pipeline and the filling order.

# Proposed approach

## State Representation



**Ensemble structure** use multiple primitives interconnected models in the main meta-model. Such pipelines also have constraints, **including primitives** and the **number of input nodes**.

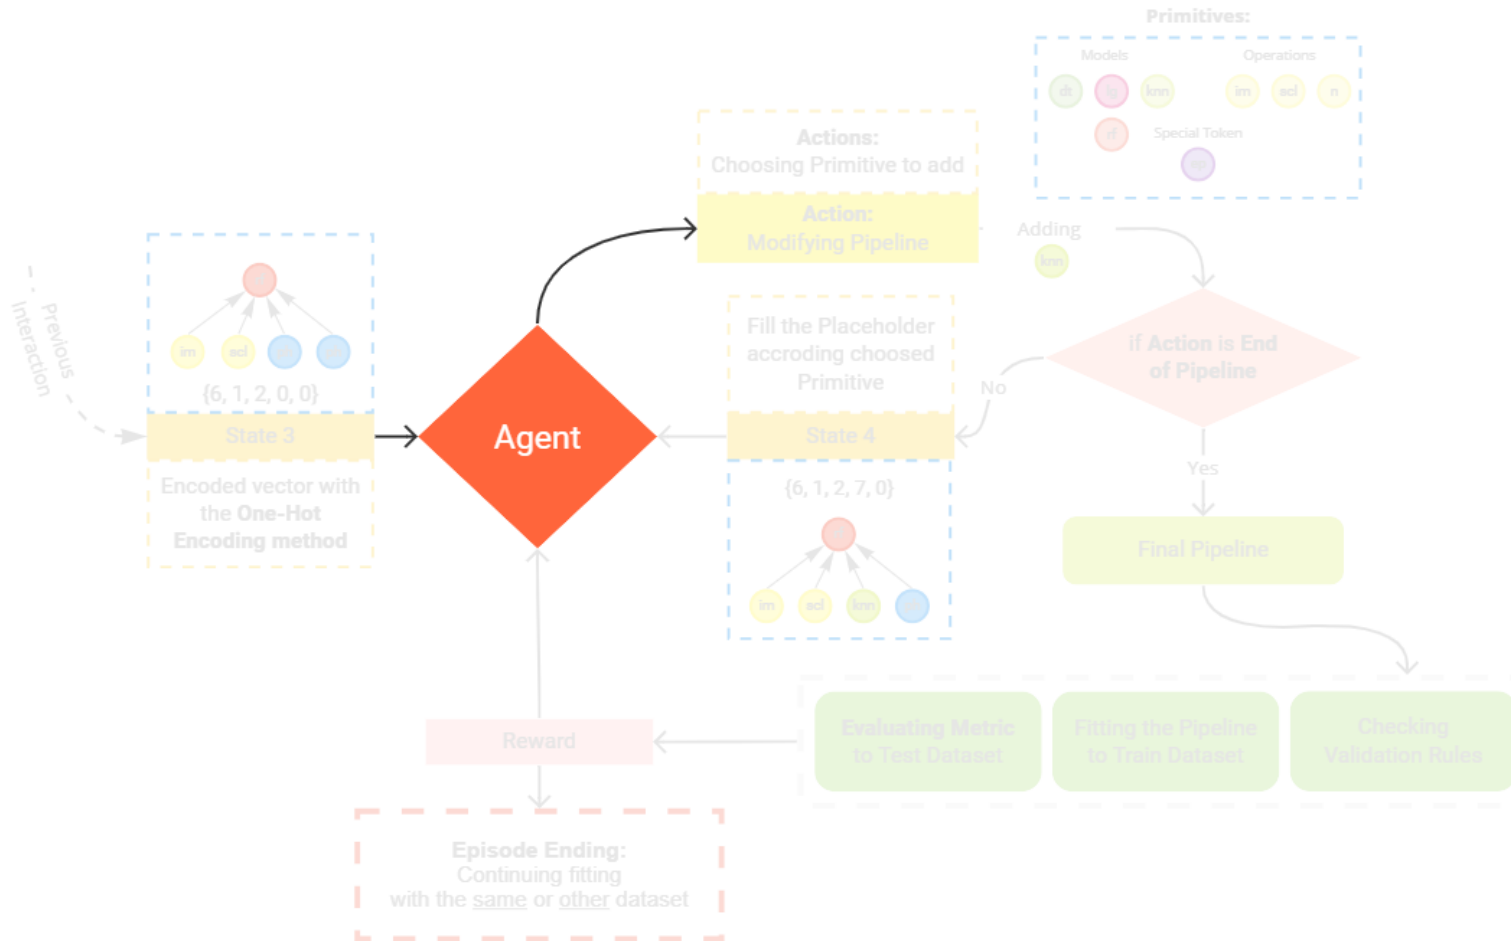
The **state** can be formalized as a **one-dimensional vector** where primitives are **encoded labels**.

Initially, the state consists of **Placeholders** and after agent's actions it fills with primitives. The filling of such ensemble pipeline starts with choosing the meta-model. Finally, the state is converted into a **pipeline**.



# Proposed approach

## Agent



The main job is done by the Agent.

The Actor-Critic algorithm (A2C) was used as the agent.

This method refers to the method based on policies.

Its purpose is to optimize the agent's behavior in various situations.

In the process, the Agent transforms the pipeline iteratively until it decides to stop.

# Proposed approach

## Agent Schema



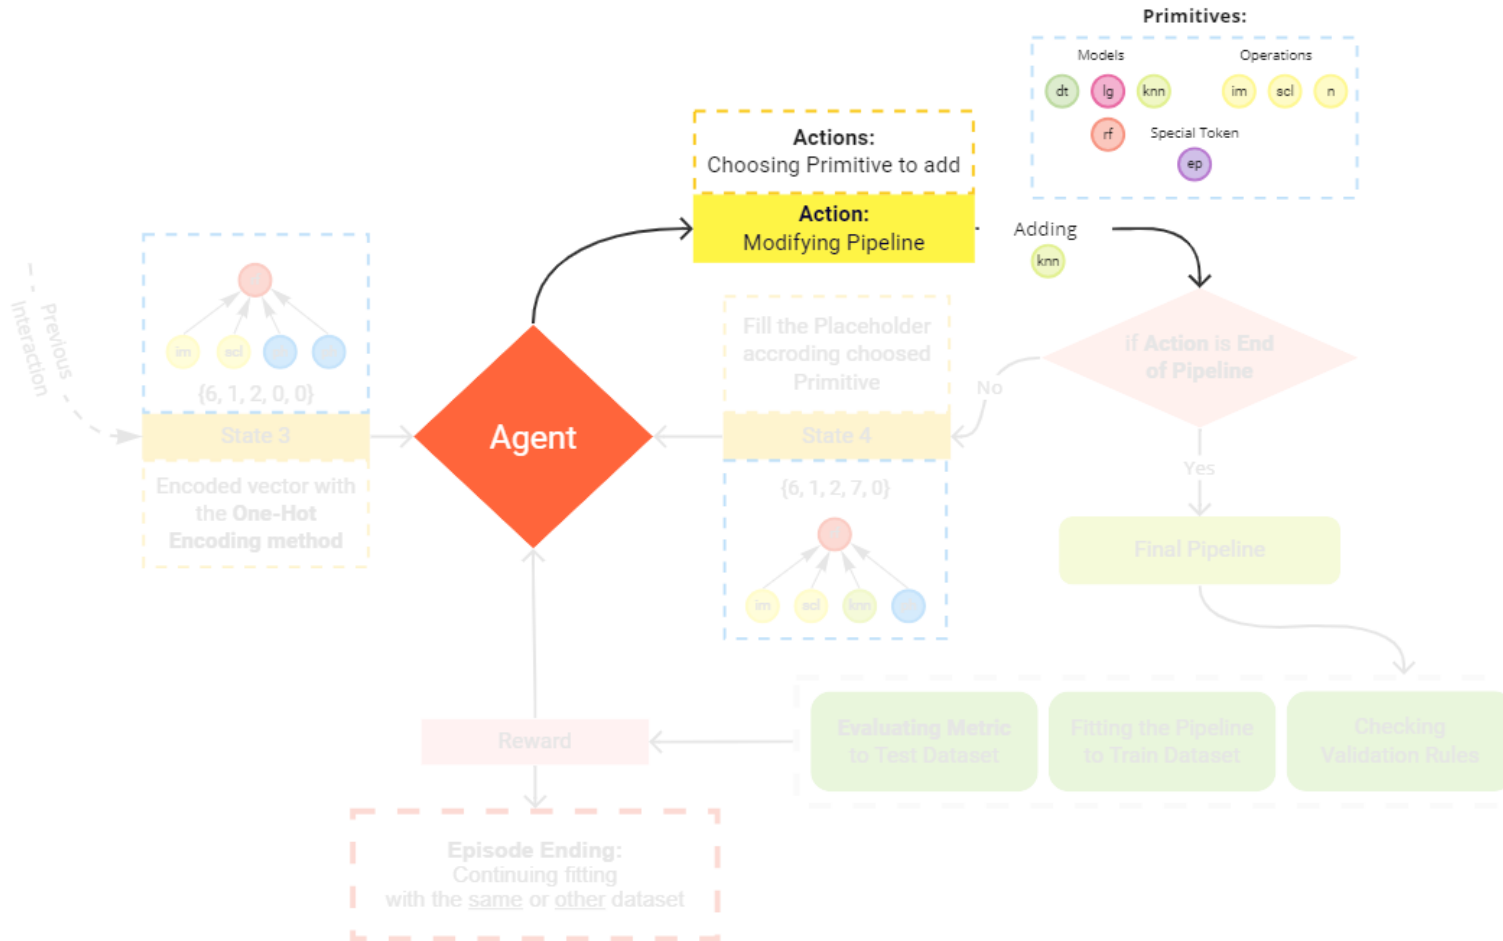
The architecture of network is consisting of 3 heads:

- Body;
- Actor;
- Critic.



# Proposed approach

## Pipeline Modifying



The **Action** is changing **Placeholder** to primitives;

Each primitive is an encoded specific **operation** or **model**.

The primitives implementation was taken from AutoML FEDOT.

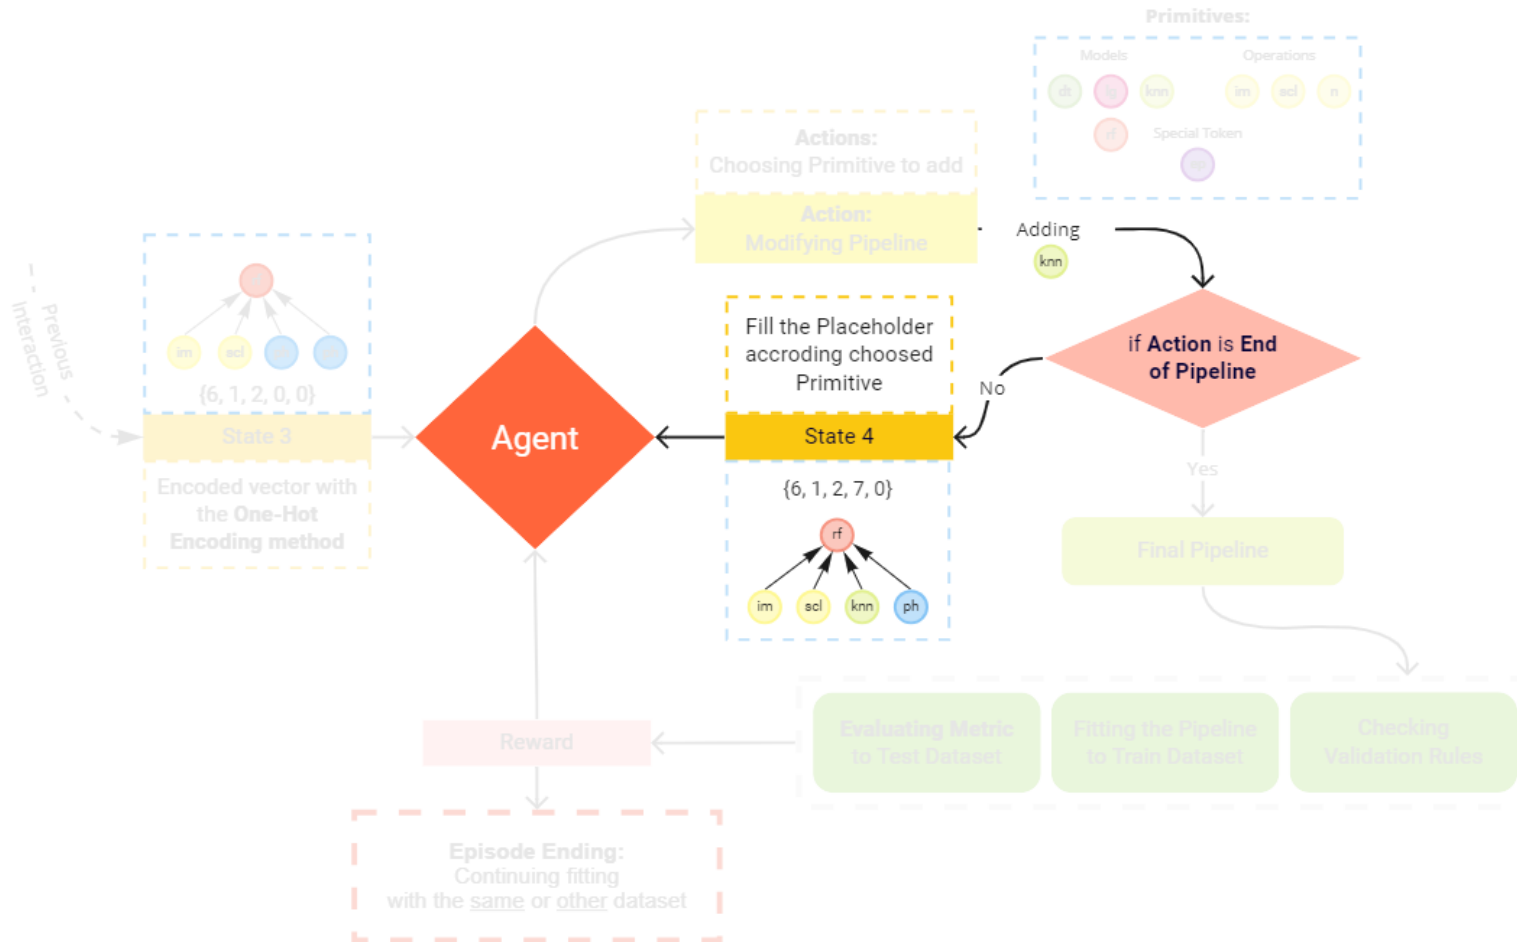
The agent can also stop the generative process early with **End of Pipeline**.

After the **Agent** chooses the action, the **Environment** checks for the ***End of Pipeline*** token.



# Proposed approach

If Not -> Continue construction

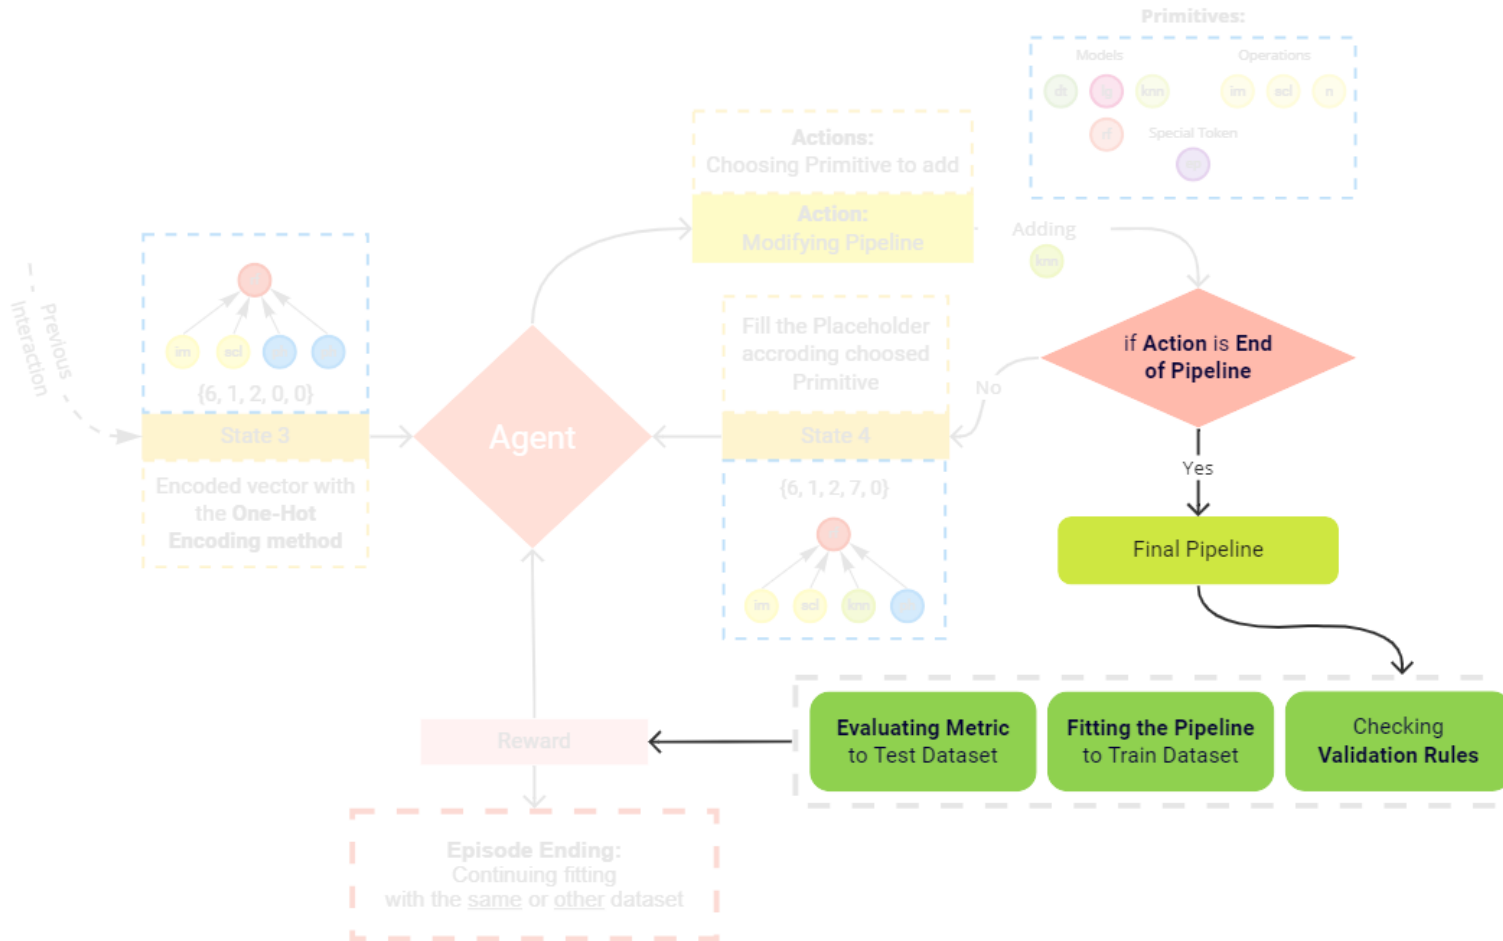


If the **Action** is not the **End of Pipeline**:

- The Environment **modifies the previous state** by chosen Operations or Models.
- The next **Placeholder** in the Pipeline is filled by chosen **primitive**.
- The **modified Pipeline**, converted into **new state**, goes to the Agent, and the cycle will continue.

# Proposed approach

If Yes -> Finish generation process

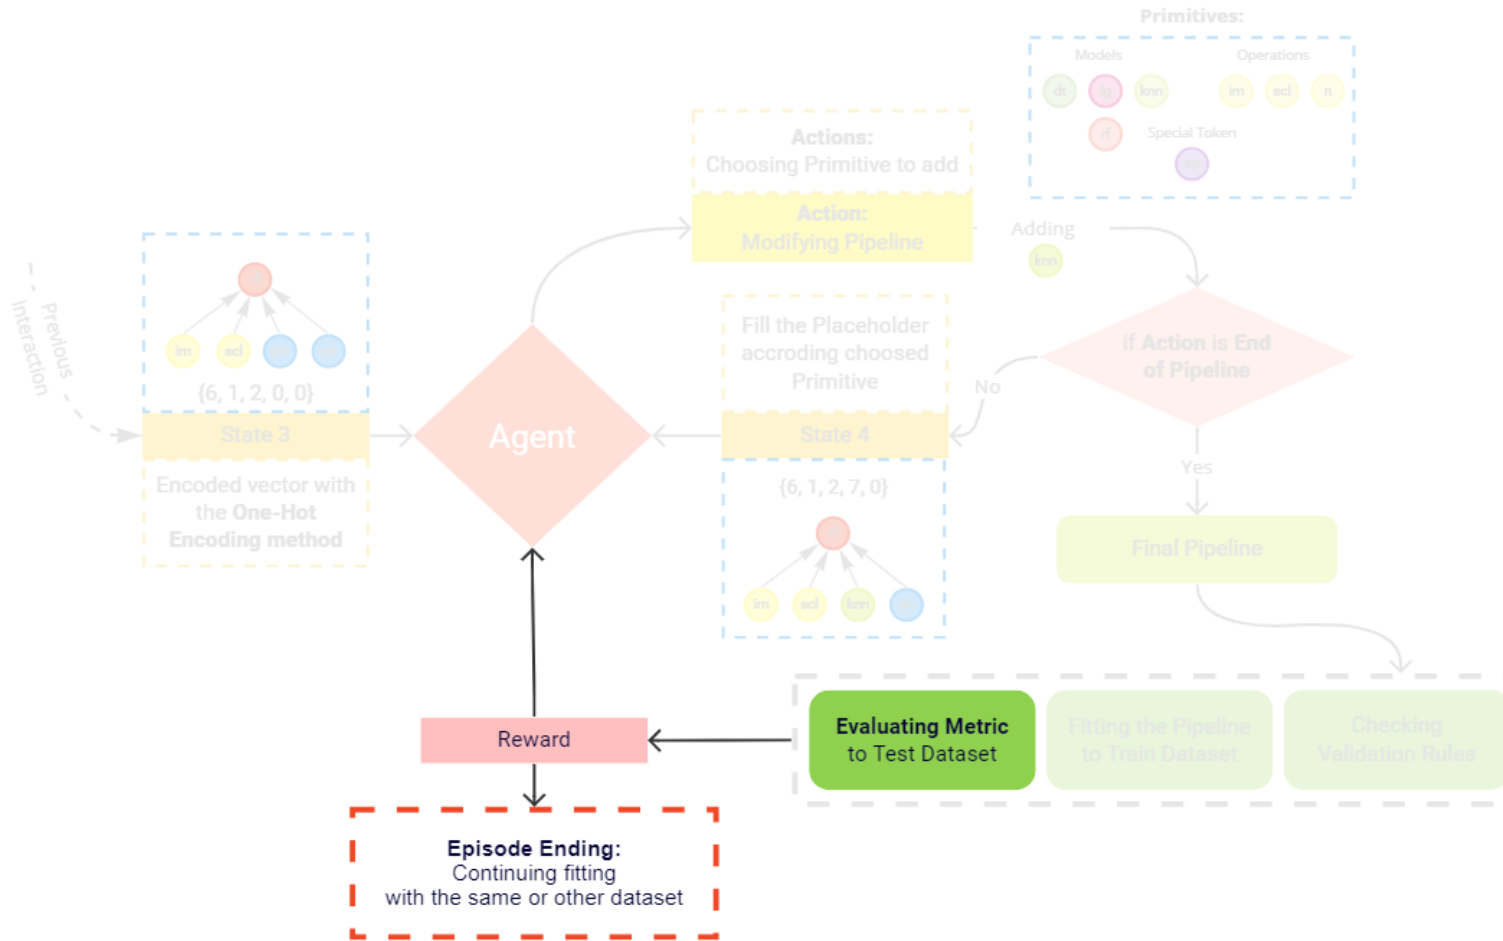


If the **Action** is the *End of Pipeline*:

- The Environment **build the final Pipeline**;
- Then, checks it at **Validation Rules**;
- **Fit** it to Train Data, if it is possible;
- Finally, **Evaluate Metric** to Test Data.

# Proposed approach

## Reward & Episode Finishing



The Agent **receives a reward** for **each interaction** with the Environment.

The **final metric score** is the most significant.

Time penalty and penalty to prevent of using a **Placeholder** or **End of Pipeline** also present in the Reward.

The proposed approach evaluated over six binary classification datasets with various specifics. Each dataset was separated into **train**, **test**, and **valid subsamples**. This was done for the **veracity of the experiment**:

- The Agent fitting at train;
- Getting reward at test;
- Validating with baselines models at valid.

Datasets	Source	Features	Columns
Amazon	Kaggle <sup>1</sup>	9	32769
Australian	OpenML <sup>2</sup>	15	690
Bank marketing	OpenML [15]	17	45211
Blood transfusion	OpenML [33]	5	748
Jasmine	OpenML <sup>3</sup>	145	2984
KC1	OpenML [24]	22	2109

The experiment was run at **Intel core i7** with parallel at **NVIDIA GeForce RTX 3060** for laptops.

The Agent was trained on **30000 episodes** on a train, and **every 1000 episodes** the Agent was verified on test data.

In the evaluation experiment, for each validation data it was required to construct **25 pipelines**. These pipelines were compared with the **baseline**.

Baseline pipelines were constructed using the **SkLearn Decision Tree model**.

Datasets	Baseline Metric	Generated Pipeline Best Metric	Generated Pipeline Mean Metric	Number of valid Pipelines
Amazon	0.6807	<b>0.8622</b>	0.8071	25
Australian	0.8439	<b>0.9505</b>	0.9140	24
Bank marketing	0.7038	<b>0.9269</b>	0.8770	22
Blood transfusion	0.6543	<b>0.7308</b>	0.7197	24
Jasmine	0.4747	0.6088	<b>0.5297</b>	24
KC1	0.5575	0.8409	<b>0.8177</b>	23

Episodes	1000		3000		10000		15000		25000		30000	
Datasets	Valid	Reward	Valid	Reward	Valid	Reward	Valid	Reward	Valid	Reward	Valid	Reward
Amazon	10	-0.42	19	0.15	25	0.68	25	0.69	25	0.64	23	0.56
Australian	10	-0.26	18	0.27	25	0.82	25	0.79	25	0.76	25	0.80
Bank	8	-0.49	19	0.25	25	0.74	25	0.77	25	0.72	25	0.79
Blood	5	-0.76	20	0.19	25	0.60	25	0.61	25	0.59	25	0.56
Jasmine	7	-0.64	22	0.16	25	0.41	25	0.41	25	0.41	25	0.42
KC1	10	-0.42	19	0.17	25	0.66	25	0.67	24	0.57	21	0.37



## Summary:

- Conducted a **research on the automatic generation of ensemble ML pipelines** using the Policy-Based RL method and presented a proof-of-concept. The proposed methods for the **environment, actions, reward, and states** were strong enough to **generate ensemble pipelines**.
- The successful generation proves the **applicability of RL methods to pipeline generation problems**.
- For future work, we plan to extend our approach to **generate non-fixed structure pipelines**.

[Source code of the experiment research](#)

Thanks for your attention